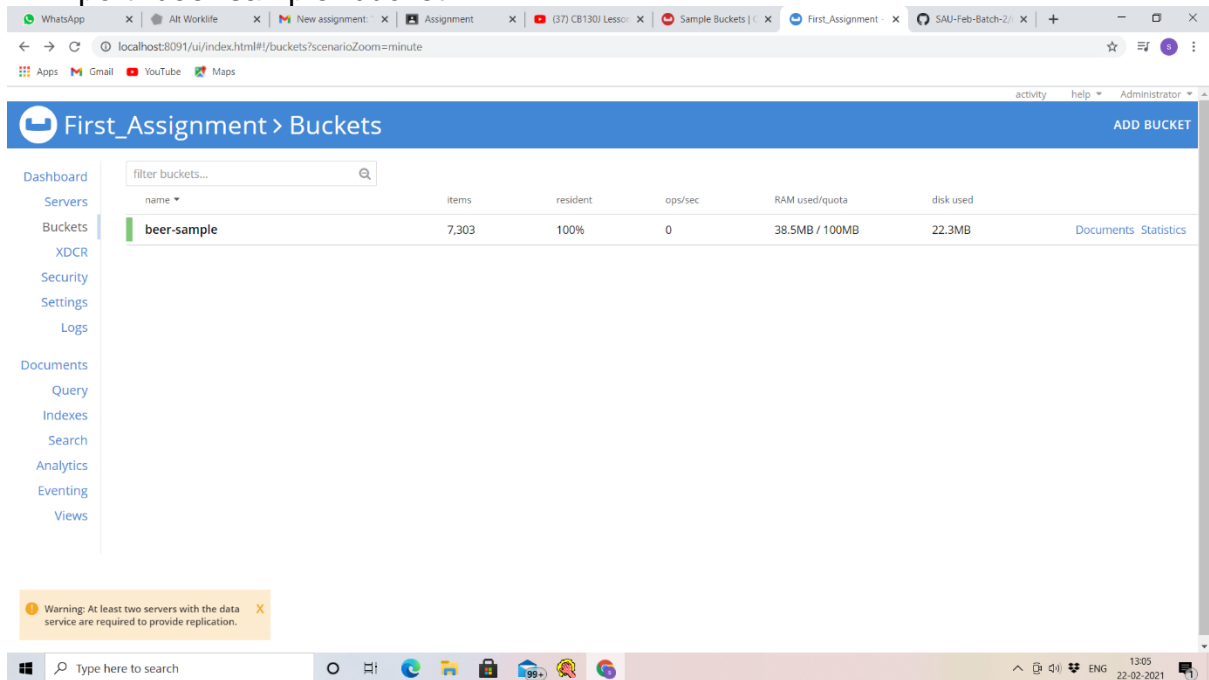


Couchbase Assignment

1. Import `beer-sample` bucket



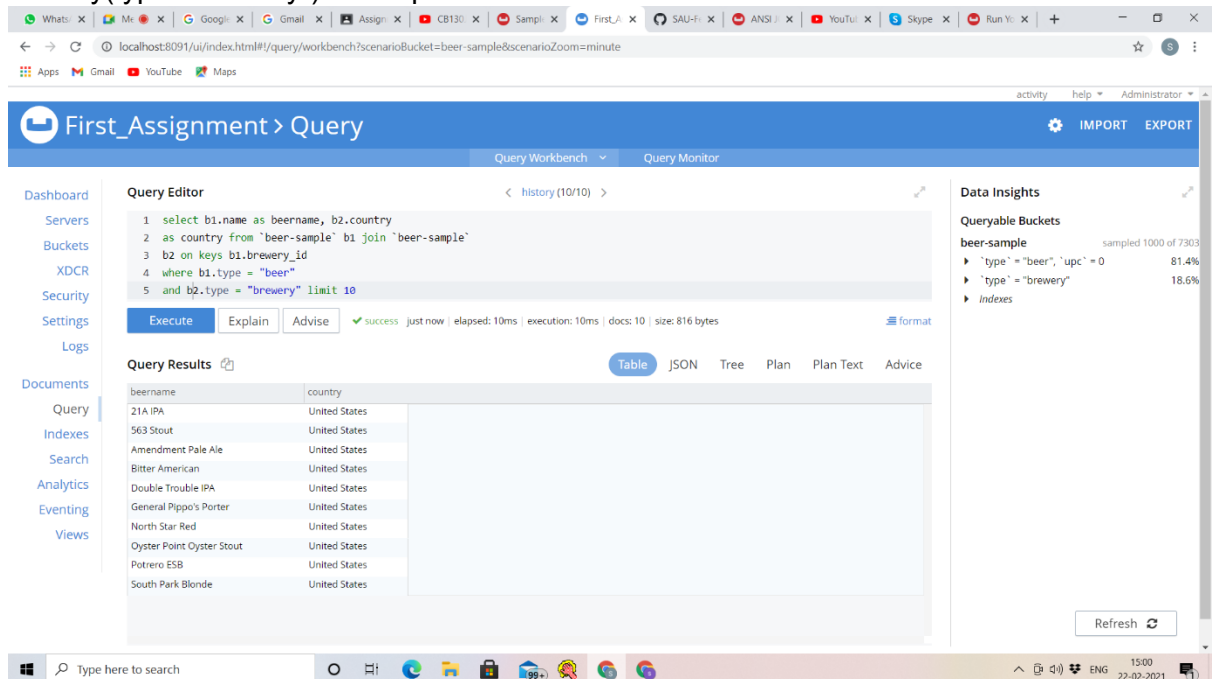
First_Assignment > Buckets

filter buckets...

name	items	resident	ops/sec	RAM used/quota	disk used
beer-sample	7,303	100%	0	38.5MB / 100MB	22.3MB

Warning: At least two servers with the data service are required to provide replication.

1. Write a join query to fetch the Top 10 brewery(type="brewery") and their country(type="beer") which produces more varieties of beers.



First_Assignment > Query

Query Editor

```
1 select b1.name as beername, b2.country
2 as country from `beer-sample` b1 join `beer-sample`
3 b2 on keys b1.brewery_id
4 where b1.type = "beer"
5 and b2.type = "brewery" limit 10
```

Execute Explain Advise success just now | elapsed: 10ms | execution: 10ms | docs: 10 | size: 816 bytes

Query Results

beername	country
21A IPA	United States
563 Stout	United States
Amendment Pale Ale	United States
Bitter American	United States
Double Trouble IPA	United States
General Pippo's Porter	United States
North Star Red	United States
Oyster Point Oyster Stout	United States
Potrero ESB	United States
South Park Blonde	United States

Data Insights

Queryable Buckets

beer-sample sampled 1000 of 7303

- ▶ 'type' = "beer", 'upc' = 0 81.4%
- ▶ 'type' = "brewery" 18.6%
- ▶ Indexes

Refresh

3. Write a mapreduce to get the number of breweries based on country. Please attach the mapreduce code and json output screenshot.

The screenshot shows a web interface for a MapReduce job. The 'Map' section contains the following code:

```
1 function (doc, meta) {  
2   if( doc.type == "brewery"){  
3     emit([doc.country, doc.type], null);  
4   }  
5 }
```

The 'Reduce' section is titled 'Reduce (built in: _count, _sum, _stats)' and contains the code:

```
1 _count
```

Below the code, the 'Results' section shows a table with the following data:

Key	Value
["", "brewery"] undefined	1
["Argentina", "brewery"] undefined	2
["Aruba", "brewery"] undefined	1

The screenshot shows the same web interface as above, but with the final results displayed in the 'Results' section. The table contains the following data:

Key	Value
["", "brewery"] undefined	1
["Argentina", "brewery"] undefined	2
["Aruba", "brewery"] undefined	1
["Australia", "brewery"] undefined	14
["Austria", "brewery"] undefined	10
["Belgium", "brewery"] undefined	99

4.

1. XDCR:

- Add a new bucket “Brewery”.
- Create a XDCR with a filter(type='brewery') to replicate only the brewery entity from `beer-sample` bucket.

The screenshot shows the 'XDCR Replications' page in the Couchbase web console. The left sidebar contains navigation links: Dashboard, Servers, Buckets, XDCR (selected), Security, Settings, Logs, Documents, Query, Indexes, Search, Analytics, Eventing, and Views. The main content area has a header 'First_Assignment > XDCR Replications' with buttons for 'ADD REPLICATION' and 'ADD REMOTE'. Below this, there are two sections: 'Remote Clusters' and 'Outgoing Replications'. The 'Remote Clusters' section shows a table with one cluster named 'First_Assignment' with IP '127.0.0.1:8091'. The 'Outgoing Replications' section shows a table with one replication from 'beer-sample' to 'Brewery' on the 'First_Assignment' cluster, with status 'replicating ..'. A warning message at the bottom states: 'Warning: At least two servers with the data service are required to provide replication.'

name	IP/hostname
First_Assignment	127.0.0.1:8091

source bucket	destination bucket	remote cluster	status
beer-sample	Brewery	First_Assignment	replicating ..

The screenshot shows the 'Buckets' page in the Couchbase web console. The left sidebar is the same as the previous screenshot. The main content area has a header 'First_Assignment > Buckets' with an 'ADD BUCKET' button. Below this is a search bar 'filter buckets...' and a table listing buckets. The table has columns: name, items, resident, ops/sec, RAM used/quota, and disk used. There are two buckets: 'beer-sample' and 'Brewery'. Both have 7,303 items and are 100% resident. 'beer-sample' has 38.9MB / 100MB RAM used and 25.3MB disk used. 'Brewery' has 34.9MB / 3.99GB RAM used and 10.1MB disk used. Links for 'Documents' and 'Statistics' are provided for each bucket. A warning message at the bottom states: 'Warning: At least two servers with the data service are required to provide replication.'

name	items	resident	ops/sec	RAM used/quota	disk used
beer-sample	7,303	100%	0	38.9MB / 100MB	25.3MB
Brewery	7,303	100%	0	34.9MB / 3.99GB	10.1MB


```
Command Prompt
-.-no-ssl-verify Skips SSL verification of certificates against CA
-.-log-file Specifies a file to log debugging information
-v,-verbose Print debugging information (Has no effect if a log
file is specified)
-t,-threads The amount of parallelism use (Default is 1)
-h,-help Prints the help message

C:\Program Files\Couchbase\Server\bin>export json -c couchbase://127.0.0.1 -u Administrator -p 9989299045 -b beer-sample -o D:/beer-sample.json -f lines -t 4
Json export failed: open D:/beer-sample.json: The system cannot find the path specified.

C:\Program Files\Couchbase\Server\bin>export json -c couchbase://127.0.0.1 -u Administrator -p 9989299045 -b beer-sample -o C:/beer-sample.json -f lines -t 4
Json export failed: open C:/beer-sample.json: Access is denied.

C:\Program Files\Couchbase\Server\bin>export json -c couchbase://127.0.0.1 -u Administrator -p 9989299045 -b beer-sample -o C:/Downloads/beer-sample.json -f lines -t 4
Json exported to "C:/Downloads/beer-sample.json" successfully
Documents exported: 7303 Documents skipped: 0

C:\Program Files\Couchbase\Server\bin>import json -c couchbase://127.0.0.1 -u Administrator -p 9989299045 -b Beer -f lines -d C:/Downloads/beer-sample.json -t 4 -g $brewery_idx
JSON import failed: file uri must be prefixed with file:///

C:\Program Files\Couchbase\Server\bin>shell export json -c couchbase://127.0.0.1 -u Administrator -p 9989299045 -b Beer -f lines -d file:///C:/Downloads/beer-sample.json -t 4 -g $brewery_idx
2021-02-24T01:37:59.657405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.658405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.660405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.662405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.663405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.664405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.664405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.664405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.665405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.666405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.666405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.667405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.667405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.669405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.672405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.674405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.674405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.675405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.675405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.676405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.676405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.677405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.678405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.678405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.679405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.679405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.680405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.681405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
2021-02-24T01:37:59.682405+30 ERROR: Error in key expression at char 0, field 'brewery_id' does not exist -- jsondata.(Parallelizer).Execute.func(1) at source.go:217
```